



Lecture 12 - Graphs, DFS, and BFS

Fall 2025, Korea University


Instructor: Gabin An (gabin_an@korea.ac.kr)

HAPPINESS IS



**... having finished
all of your exams!**

Midterm Exam



- Average: 61 (std: 20.5)
- Median: 59
- Please come see me after class if you'd like to check your exam paper.

Course Outline (After Midterm)

- Part 3: Data Structures - Continued
 - **Graphs, Graph Search and Applications** ↗
- Part 4: Dynamic Programming
 - Shortest-Path: Dijkstra, Bellman-Ford, Floyd-Warshall Algorithms
 - More General DP: Longest Common Subsequence, Knapsack Problem
- Part 5: Greedy Algorithms and Others
 - Scheduling Problem, Optimal Codes
 - Minimum Spanning Trees
 - Max Flow, Min Cut and Ford-Fulkerson Algorithms
 - Stable Matching, Gale-Shapley Algorithm


Today's Objectives

- Learn the basics of **graphs**, including terminology like **vertices**, **edges**, and types of graphs.
- Understand different methods for representing graphs in memory: **Adjacency Matrix** and **Adjacency List**.
- Explore **Depth First Search (DFS)**
- Explore **Breadth First Search (BFS)**
- Compare and contrast DFS and BFS to understand their core differences.

What are Graphs?

- A graph $G = (V, E)$ is a set of vertices (nodes) V and edges E connecting them.
- n = number of vertices ($|V|$), m = number of edges ($|E|$).
- **Directed vs. Undirected:**
 - **Undirected:** edges are bidirectional. $(i, j) \in E \iff (j, i) \in E$.
 - **Directed:** edges have a direction.
- **Connected Graph:** A path exists between any two nodes.
- **Sparse vs. Dense:**
 - **Sparse:** Few edges, e.g., $m = \Theta(n)$.
 - **Dense:** Many edges, e.g., $m = \Theta(n^2)$.


Example: Directed Graph w/ Six Nodes and Six Edges



Degree of a Node

- The number of edges connected to that node.
- For undirected graphs:
 - $\deg(u)$ = number of edges incident to u
- For directed graphs:
 - Each node has two types of degrees:
 - In-degree: number of incoming edges (edges ending at the node)
 - Out-degree: number of outgoing edges (edges starting from the node)


Representing Graphs




- How to store a graph in memory? Two main methods:
 - i. **Adjacency Matrix**
 - ii. **Adjacency List**

Representing Graphs - 1. Adjacency Matrix

- An $n \times n$ binary matrix.
- Matrix element (i, j) is **1** if edge (i, j) exists, **0** otherwise.




Representing Graphs - 1. Adjacency Matrix - Continued




- **Space:** $\Theta(n^2)$ (why? $n \times n$).
- **Edge Lookup:** $O(1)$ time to check if an edge exists.
- **Neighbor Enumeration:** $\Theta(n)$ time.

Representing Graphs - 2. Adjacency List



- An array A of n lists.
- $A[u]$ contains a list of all vertices v where $(u, v) \in E$.




Array Linked List



Representing Graphs - 2. Adjacency List - Continued



Array **Linked List**



- **Space:** $\Theta(n + m)$.
- **Edge Lookup:** $O(\deg(u))$ time.
- **Neighbor Enumeration:** $O(\deg(u))$ time.

Representing Graphs - Summary

Representation	Space Complexity	Check if edge (u, v) exists	List all neighbors of u	Notes
Adjacency Matrix	$\Theta(n^2)$	$\Theta(1)$	$\Theta(n)$	Fast edge lookup but uses lots of memory.
Adjacency List	$\Theta(n + m)$	$O(\deg(u))$ ($\approx O(V)$ worst case)	$\Theta(\deg(u))$	Efficient for sparse graphs.

Graph Traversal Algorithms

- Depth-First Search (DFS)
- Breadth-First Search (BFS)

Depth First Search (DFS)


- **Strategy:** Explores a path **as far as possible** before backtracking.
- **DFS Augmentations (to prevent loops):**
 - **Colors:**
 - **White:** Unvisited.
 - **Gray:** Visited, but not all neighbors explored.
 - **Green :** Completely processed.
 - **Parent Pointers:** Keeps track of the traversal path.
 - **Timestamps:**
 - **Discovery time ($d(v)$):** When a node is first visited.
 - **Finish time ($f(v)$):** When a node is completely processed.

DFS Pseudocode


- Initially, every node is marked as unvisited (white).

```
Algorithm DFS(s, t):
    color(s) ← gray
    d(s) ← t
    t ← t + 1
    foreach v in N(s) do
        if color(v) == white then
            p(v) ← s
            t ← DFS(v, t)
            t ← t + 1
    f(s) ← t
    color(s) ← green
    return f(s)
```


DFS Example




- Start at a white node **a**.




- From node **a**, we will visit all of **a**'s children, namely node **b**.




- We now visit **b**'s child, node **c**.





- We will next search `c`'s second child, node `d`, as another child, node `a`, has already been visited (i.e., colored gray).




- Since node `d` has no children, we return back to its parent node, `c`, and continue to go back up the path we took, marking nodes with a finish time when we have searched all of their children.











- We start with a new source node `e` and run DFS to completion.




DFS Runtime

- **Runtime:** $O(\# \text{ of node visits} + \# \text{ of edge scans})$.
- Each node is visited at most once, so node visits are $O(n)$.
- Each edge is scanned at most twice (or once for directed graphs), so edge scans are $O(m)$.
- **Total Runtime:** $O(m + n)$.

DFS and Stack




Stack (Last-In-First-Out; LIFO)



Breadth First Search (BFS)

- **Strategy:** Expands the search frontier uniformly, visiting all nodes at distance k before moving to nodes at distance $k + 1$.
- **Data Structure:** A **queue** (First-In, First-Out) is used to explore nodes in order of their distance from the source.
- **Level Sets (L_i):** BFS computes sets L_i containing all nodes at distance i from the source.


Level Sets?




BFS Pseudocode

```
Algorithm BFS(s):
    vis[v] ← false for all v
    L_0 ← {s}
    vis[s] ← true
    for i = 0...n-1 do
        if L_i is empty then exit
        while L_i is not empty do
            u ← L_i.pop()
            foreach x in N(u) do
                if vis[x] == false then
                    vis[x] ← true
                    L_{i+1}.insert(x)
                    p(x) ← u
```


BFS Example




- We begin with the first level set, L_0 , which contains only the source node a .




- Next, we explore the second level set, L_1 , consisting of nodes directly reachable from L_0 .



- We then move on to the third level set, L_2 , which contains nodes directly reachable from L_1 .




- Finally, we reach the last level set, L_4 . With no further children to explore, the search terminates here.



BFS Runtime


- **Runtime:** $O(\# \text{ nodes visited} + \# \text{ edges scanned})$.
- Each node is visited once, so node visits are $O(n)$.
- Each edge is scanned once, so edge scans are $O(m)$.
- **Total Runtime:** $O(m + n)$.

BFS and Queue



BFS

Queue (First-In-First-Out; FIFO)




BFS vs. DFS Simplified

- **DFS:** Explore as deep as possible before backtracking
 - Uses a **stack** to go as far down a path as possible before backtracking. It's like exploring a single winding tunnel to its end before coming back.
- **BFS:** Explore all neighbors level by level
 - Uses a **queue** to explore nodes closest to the source first. It's like ripples spreading out in a pond.




BFS vs. DFS Simplified




DFS

BFS

Stack (Last-In-First-Out; LIFO)



Queue (First-In-First-Out; FIFO)



Credits & Resources

Lecture materials adapted from:

- Stanford CS161 slides and lecture notes
 - <https://stanford-cs161.github.io/winter2025/>
- *Algorithms Illuminated* by Tim Roughgarden
 - <https://algorithmsilluminated.com/>